

Mysticeti Consensus Digital Twin

Amp + Wolfram Mathematica + Lean 4 + Rust

An executable + formally bounded microscope for the quorum and direct-decision core behind production Sui Layer-1 consensus.

Scientific boundary Bounded Rust research twin · not production equivalence · not a production benchmark reproduction



Why Mysticeti matters

A research result with a real deployment path — and an unusually sharp latency target.

Oct 2023

Preprint



“Reaching the Latency Limits with Uncertified DAGs”

25 Jul 2024

Sui Mainnet



Validators switched to Mysticeti-C consensus.

NDSS 2025

Peer-reviewed



Mysticeti appears in the NDSS 2025 program.

Goal

3 message rounds



Direct commit path targets three message rounds.

SOURCE TRAIL arxiv.org/abs/2310.14821 · ndss-symposium.org/ndss-paper/mysticeti · docs.sui.io/.../consensus

Production relevance raises the assurance bar: every claim needs a named evidence lane.

One protocol. Five different kinds of claim.

The challenge is not generating output — it is preventing evidence from silently changing meaning.

PAPER CLAIM**What the authors state**

citation + normalized claim

TRACEABLE**RUST DIGITAL TWIN****What bounded traces execute**

event-driven replay + audit

TRACEABLE**FORMAL PROOF****What follows from assumptions**

Lean kernel-checked theorem

TRACEABLE**REPORTED DATA****What Table I reports**

local CSV transcription


TRACEABLE**PRODUCTION SUI****What deployed Sui does**

public sources; parity not claimed

OUT OF SCOPE**No upward inference without a bridge.**

Amp coordinates three lanes — one evidence gate

Rust executes traces, Mathematica explores and audits, Lean



EVIDENCE GATE Every output carries: source → assumption → method → result → limitation.

Evidence hierarchy — strong does not mean broad

Confidence is attached to the exact statement, not inherited by neighboring protocol claims.

FORMALLY PROVED

Lean checks the theorem from explicit finite-set assumptions.

**STRONG +
NARROW****RUST EXECUTABLE**

Bounded stake-weighted event traces, strict replay and audit.

MATHEMATICA

Independent conformance, fixtures and visualization.

CSV TRANSCRIPTION

Locally stored paper values — not independent reproduction.


OUT OF SCOPE

Production parity · full protocol proof · liveness · performance.

**NO
CLAIM**

Read the DAG as evidence — not as decoration

Rounds run left-to-right in the source graphic; validator lanes and authority identity govern every count.



ROUNDS Causal depth, not wall-clock time

AUTHORITY One validator = one voting identity

EQUIVOCATION More blocks $\neq$ more voting power

PARENTS Distinct-authority references

The generated fixture explicitly checks parent references and distinct parent authorities.

A certificate is not a commit


The model separates evidence at r , $r+1$, and $r+2$ so “I saw support” cannot become accidental finality.



DIRECT SKIP At $r + 1$, q distinct non-support authorities can justify a conservative direct skip — no $r + 2$ certificate quorum required.

Consensus Safety Microscope

A developer changes assumptions, predicts the outcome, then receives inspectable evidence — not a magic score.



fault bound f




crashed validators



round count



equivocation



random seed



RETURNED EVIDENCE

Decision Commit / Skip / Undecided

Support authority IDs at $r + 1$

Certificates block + authority IDs at $r + 2$

Threshold explicit $q = 2f + 1$

Invariants parents valid · authorities distinct

PREDICT → VARY → INSPECT

The executable twin is small enough to inspect — complete enough to challenge

Every transition is deterministic; every recorded claim can be replayed and independently audited.

COMMITTEE

validated, stake-weighted

CANONICAL DAG

transactional insertion

LOCAL VIEWS

one per authority

EVENT QUEUE

deterministic ordering

DECISION ENGINE

direct commit / skip

CAMPAIGN ENGINE

dedicated Rayon pool

STRICT REPLAY

schema + tamper checks

INVARIANT AUDITOR

independent evidence checks

TRACE CONTRACT

inputs → receiver-specific events → local DAG transitions → decisions → invariant
evidence → strict replay

RESEARCH TWIN ≠ PRODUCTION SUI

Faults matter because delivery changes what each authority can know

The simulator models causal consequences, not just labels attached after a global DAG is built.



The release evidence is executable, not decorative

Actual Windows + Ubuntu WSL reports: tests, clippy, trace replay, deterministic campaigns, and Criterion evidence.

```
● ● ● WSL / Linux reproducibility evidence

$ cat exports/wsl_rust_build_report.txt
Ubuntu 24.04.3 LTS | WSL2 | 24 vCPUs
Rust test result: 26 passed; 0 failed; 0 ignored
cargo clippy --all-targets --all-features -- -D warnings PASS

TRACE Windows == Ubuntu WSL
6272fa854de66bc42512f38d095d7ccf6f75bb85f581dc4acaabf9fbe8ede71d
BYTE_IDENTICAL_WINDOWS_WSL_TRACE=true

CAMPAIGN Win jobs=1 == Win jobs=8 == WSL jobs=1 == WSL jobs=8
20 seeds / 80 rows
5fdc665f5c35cee5c63143860789a9a2a4db831ec65c9e814610d1f6b1764a6a
BYTE_IDENTICAL_WINDOWS_WSL_JOBS1_JOBS8=true

$ cat exports/wsl_benchmark_report.txt
Criterion 32 authorities / 50 slots | 10 samples
time: [2.9226 s, 2.9977 s estimate, 3.0633 s]
Development-machine WSL evidence; not bare-metal or production.
```

26 TESTS · 0 FAILED / IGNORED

PARALLEL CAMPAIGN TESTS

CLIPPY: WINDOWS + WSL

WINDOWS [3.5361, 3.5774] s

WSL EST. 2.9977 s

Same evidence across OS and scheduler

Independent scenario concurrency changes timing and interleaving — never canonical evidence ordering.




BOUNDARY Dedicated Rayon pool across independent scenarios · each scenario deterministic · canonical sort before CSV serialization · no global pool mutation.

WSL CRITERION 32 authorities / 50 slots · [2.9226 s, 2.9977 s estimate, 3.0633 s] · 10 samples · Ubuntu WSL2 · 24 vCPUs

NON-CLAIMS Development-machine WSL evidence; not bare-metal/production; no speedup comparison against Windows; no Ethereum builder flow or Sui production implementation.

Loss reduces progress in the bounded Rust campaign

20 seeds × 4 loss levels = 80 rows; dedicated Rayon pool across independent deterministic scenarios.



LOSS 0.00

8.0 average commits

LOSS 0.30


2.7 average commits

DETERMINISM BOUNDARY

**Canonical sort before CSV;
no global pool mutation.**

A reusable Wolfram architecture — not a notebook monolith

Deterministic package APIs feed notebook exploration, fixtures, plots, CSV evidence, and release reports.



```
SupportedProposal[dag_, block_, slot_] :=  
breadth-first ancestors → sort {round, validator, ID}  
→ return at most one proposal deterministically
```

**Authority deduplication is inside the evidence path —
equivocation never manufactures quorum weight.**

The quorum proof chain, end to end

The kernel proves finite counting under equal voting power — then exposes the honest-node obligation.



At most f Byzantine intersection has at least $f + 1$ $\Rightarrow \geq 1$ **honest validator**

GLOBAL COUNTING reduces safety to a local invariant: **an honest validator does not double-support.**

Exact Lean coverage — theorem by theorem

Names are part of the audit trail. Green states what is proved; the right edge states what is not.

`quorum_intersection_at_least_f_add_one`

Two $2f+1$ quorums inside $3f+1$ overlap by $\geq f+1$.

`quorum_intersection_contains_honest`

With $\leq f$ Byzantine, the intersection contains an honest validator.

`lemma5_at_most_one_certified_block_per_slot`

Same-slot certificates coincide, given honest single-support.

`conflicting_transactions_cannot_both_gather_quorums` Conflicting transactions cannot both gather quorums, given honest no-double-vote.

DOES NOT PROVE

DAG traversal
Direct + indirect decisions
Networking + signatures
Liveness + epoch change
Rust equivalence


PROOF HYGIENE

PASSED — no sorry • admit • axiom • unsafe • native_decide

LEAN BUILD: 3001 JOBS

Safety envelope: where honest overlap survives

This chart is exact for the equal-authority abstraction; the margin is $\max(0, 2q - n - b)$.



READ IT

Positive margin $\Rightarrow$ at least one honest validator remains in quorum overlap.

BOUNDARY

Zero margin $\Rightarrow$ this counting argument no longer guarantees honest overlap.


LIMITATION

Production Sui uses stake-weighted voting power. This figure counts equal authorities.

EQUAL AUTHORITY ONLY

Fault sensitivity: useful shape, deliberately uncalibrated

The model asks “how does the response move?” — never “what will production latency be?”




MODEL INPUTS

3 × base delay
+ Gaussian jitter
+ 45 ms / crash
+ 25 ms / Byzantine
≥ 250 seeded samples

**NOT A WAN SIMULATION
NOT A PRODUCTION PREDICTION**

Production comparison: faithful transcription, not reproduction

NDSS 2025 Table I · committee 137 · offered load 5,000 TPS. Values are checked against a local CSV fixture.



BULLSHARK

P50 2890 ms
P95 4600 ms

MYSTICETI-C

P50 650 ms
P95 975 ms

CLAIM LIMIT

CSV consistency only. No independent deployment or measurement.

What developers should take into code review

The theorem is small. The implementation obligations it exposes are concrete.

01

COUNT AUTHORITY

Never count blocks or messages as independent voting power.

02

SEPARATE AXES

Crash-performance sensitivity is not Byzantine safety.

03

LOCAL OBLIGATION

Persist “no double-support / no conflicting vote” across recovery.

04

WEIGHTED REALITY

Production Sui uses delegated-stake voting power > 2/3.

05

TRACE ADAPTER PATH

Builder/relay-style timestamp + order traces → provenance → revision-pinned differential tests. No Ethereum builder flow or Sui production implementation claim.

One command. Three lanes. One bounded release gate.

Rust executes and replays; Wolfram conforms and visualizes; Lean checks the mapped theorem.

```
PS> Set-Location "C:\Amp_demos\Mysticeti-Consensus-Digital-Twin"; .\build_all.ps1
```

RUST

26 tests · clippy × 2 OS

CAMPAIGN

jobs1 = jobs8 hash

LEAN 4

3001 jobs · scan pass

EVIDENCE GATE

audit 29/0 · PASSED

26

Rust tests

5fdc...a6a

jobs1 = jobs8

12 / 12

Wolfram

29 / 0

audit pass/fail


PASS

Lean proof hygiene

COMBINED RELEASE GATE: PASSED


From one digital twin to a reusable consensus research workbench

Proposed 16-week plan: reuse the experiment machinery, but keep every protocol's state machine, assumptions, and proofs explicit.



A four-month path to a public multi-protocol beta

Sequence by learning value: establish trusted CFT baselines, then classic/chained BFT, then the DAG stack and cross-protocol evidence.



PUBLIC-BETA GATE

deterministic replay · comparable metrics · explicit assumptions · protocol-specific theorem coverage · reproducible examples

Planning principle: shared infrastructure, never a “universal proof.” Every adapter ships with its own semantics, invariants, exclusions, and evidence.

LET'S BUILD THE NEXT BRIDGE

Bring proofs to production conversations.

Rust executes and replays. Wolfram explores and audits. Lean proves bounded mathematics. Amp orchestrates the evidence gate.

Daniel Liezrowice

ESL · Engineering Software Lab



<https://eswlab.com/>

<https://il.linkedin.com/in/liezrowice>

<https://github.com/zuwasi/mysticeti-proof-to-production-demo>

Independent bounded research demo; not affiliated with or endorsed by Mysten Labs or Sui. Lean proves mapped mathematics, not the Rust binary.

